

class09_mini_project

Hyeonseok Jang (PID# A59011126)

10/27/2021

1. Exploratory data analysis

```
# Save your input data file into your Project directory  
fna.data <- "WisconsinCancer.csv"
```

```
# Complete the following code to input the data and store as wisc.df  
wisc.df <- read.csv(fna.data, row.names=1)
```

```
# We can use -1 here to remove the first column  
wisc.data <- wisc.df[,-1]
```

```
# Create diagnosis vector for later  
diagnosis <- as.factor(wisc.df$diagnosis)
```

Q1. How many observations are in this dataset?

```
nrow(wisc.data)
```

```
## [1] 569
```

There are 569 observations in the dataset.

Q2. How many of the observations have a malignant diagnosis?

```
table(diagnosis)
```

```
## diagnosis  
##    B    M  
## 357 212
```

212 observations have a malignant diagnosis.

Q3. How many variables/features in the data are suffixed with `_mean`?

```
length(grep("_mean", colnames(wisc.df)))
```

```
## [1] 10
```

There are 10 variables in the data with the '_mean' suffix.

2. Principal Component Analysis Performing PCA

```
# Check column means and standard deviations
colMeans(wisc.data)
```

##	radius_mean	texture_mean	perimeter_mean
##	1.412729e+01	1.928965e+01	9.196903e+01
##	area_mean	smoothness_mean	compactness_mean
##	6.548891e+02	9.636028e-02	1.043410e-01
##	concavity_mean	concave.points_mean	symmetry_mean
##	8.879932e-02	4.891915e-02	1.811619e-01
##	fractal_dimension_mean	radius_se	texture_se
##	6.279761e-02	4.051721e-01	1.216853e+00
##	perimeter_se	area_se	smoothness_se
##	2.866059e+00	4.033708e+01	7.040979e-03
##	compactness_se	concavity_se	concave.points_se
##	2.547814e-02	3.189372e-02	1.179614e-02
##	symmetry_se	fractal_dimension_se	radius_worst
##	2.054230e-02	3.794904e-03	1.626919e+01
##	texture_worst	perimeter_worst	area_worst
##	2.567722e+01	1.072612e+02	8.805831e+02
##	smoothness_worst	compactness_worst	concavity_worst
##	1.323686e-01	2.542650e-01	2.721885e-01
##	concave.points_worst	symmetry_worst	fractal_dimension_worst
##	1.146062e-01	2.900756e-01	8.394582e-02

```
apply(wisc.data,2,sd)
```

##	radius_mean	texture_mean	perimeter_mean
##	3.524049e+00	4.301036e+00	2.429898e+01
##	area_mean	smoothness_mean	compactness_mean
##	3.519141e+02	1.406413e-02	5.281276e-02
##	concavity_mean	concave.points_mean	symmetry_mean
##	7.971981e-02	3.880284e-02	2.741428e-02
##	fractal_dimension_mean	radius_se	texture_se
##	7.060363e-03	2.773127e-01	5.516484e-01
##	perimeter_se	area_se	smoothness_se
##	2.021855e+00	4.549101e+01	3.002518e-03
##	compactness_se	concavity_se	concave.points_se
##	1.790818e-02	3.018606e-02	6.170285e-03
##	symmetry_se	fractal_dimension_se	radius_worst
##	8.266372e-03	2.646071e-03	4.833242e+00
##	texture_worst	perimeter_worst	area_worst
##	6.146258e+00	3.360254e+01	5.693570e+02

```
##      smoothness_worst      compactness_worst      concavity_worst
##      2.283243e-02      1.573365e-01      2.086243e-01
##      concave.points_worst      symmetry_worst      fractal_dimension_worst
##      6.573234e-02      6.186747e-02      1.806127e-02
```

```
# Perform PCA on wisc.data by completing the following code
wisc.pr <- prcomp(wisc.data, scale=TRUE)
```

```
# Look at summary of results
summary(wisc.pr)
```

```
## Importance of components:
##              PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation    3.6444 2.3857 1.67867 1.40735 1.28403 1.09880 0.82172
## Proportion of Variance 0.4427 0.1897 0.09393 0.06602 0.05496 0.04025 0.02251
## Cumulative Proportion 0.4427 0.6324 0.72636 0.79239 0.84734 0.88759 0.91010
##              PC8      PC9      PC10     PC11     PC12     PC13     PC14
## Standard deviation    0.69037 0.6457 0.59219 0.5421 0.51104 0.49128 0.39624
## Proportion of Variance 0.01589 0.0139 0.01169 0.0098 0.00871 0.00805 0.00523
## Cumulative Proportion 0.92598 0.9399 0.95157 0.9614 0.97007 0.97812 0.98335
##              PC15     PC16     PC17     PC18     PC19     PC20     PC21
## Standard deviation    0.30681 0.28260 0.24372 0.22939 0.22244 0.17652 0.1731
## Proportion of Variance 0.00314 0.00266 0.00198 0.00175 0.00165 0.00104 0.0010
## Cumulative Proportion 0.98649 0.98915 0.99113 0.99288 0.99453 0.99557 0.9966
##              PC22     PC23     PC24     PC25     PC26     PC27     PC28
## Standard deviation    0.16565 0.15602 0.1344 0.12442 0.09043 0.08307 0.03987
## Proportion of Variance 0.00091 0.00081 0.0006 0.00052 0.00027 0.00023 0.00005
## Cumulative Proportion 0.99749 0.99830 0.9989 0.99942 0.99969 0.99992 0.99997
##              PC29     PC30
## Standard deviation    0.02736 0.01153
## Proportion of Variance 0.00002 0.00000
## Cumulative Proportion 1.00000 1.00000
```

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

44.27% of the original variance is captured by the PC1.

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?


Three PCs

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

seven PCs

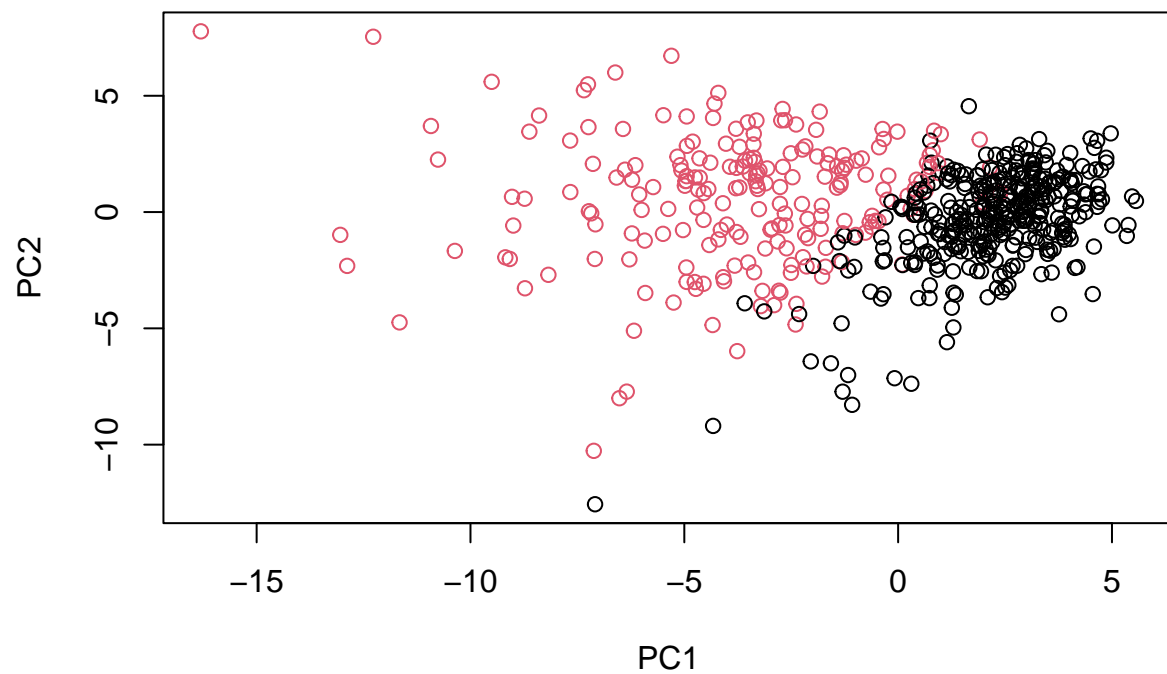
Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

```
biplot(wisc.pr)
```



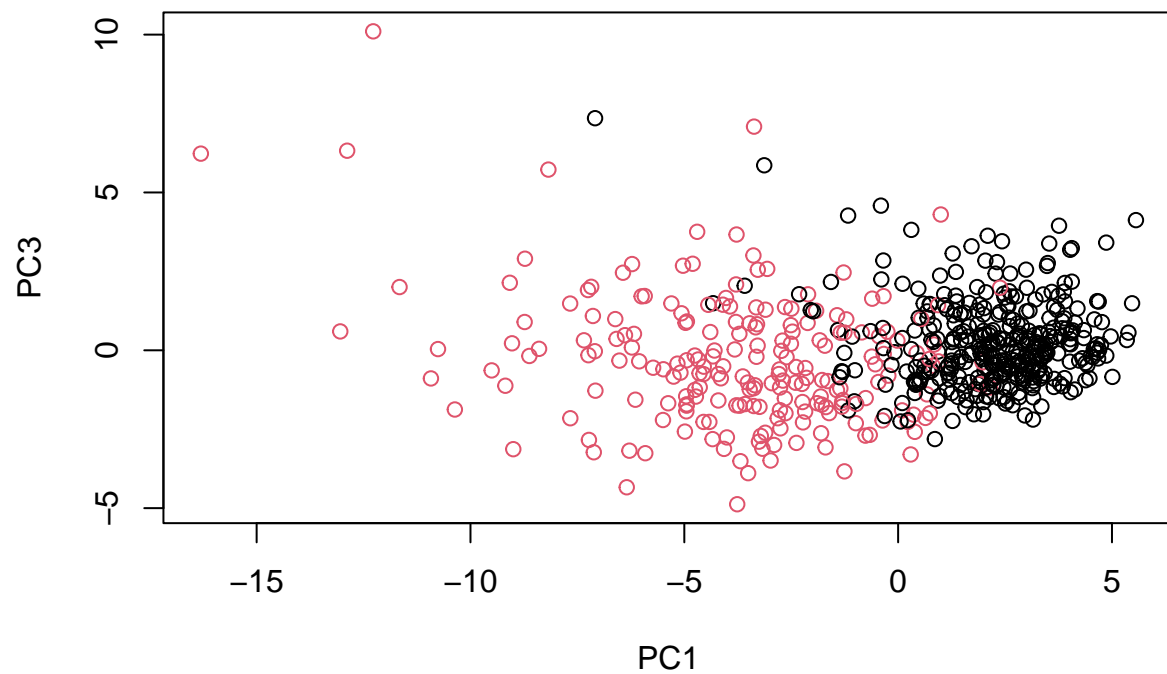
Rownames are used as the plotting character for biplots like this one which can make trends rather hard to see.

```
# Scatter plot observations by components 1 and 2
plot(wisc.pr$x[,1:2], col = diagnosis, xlab = "PC1", ylab = "PC2")
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

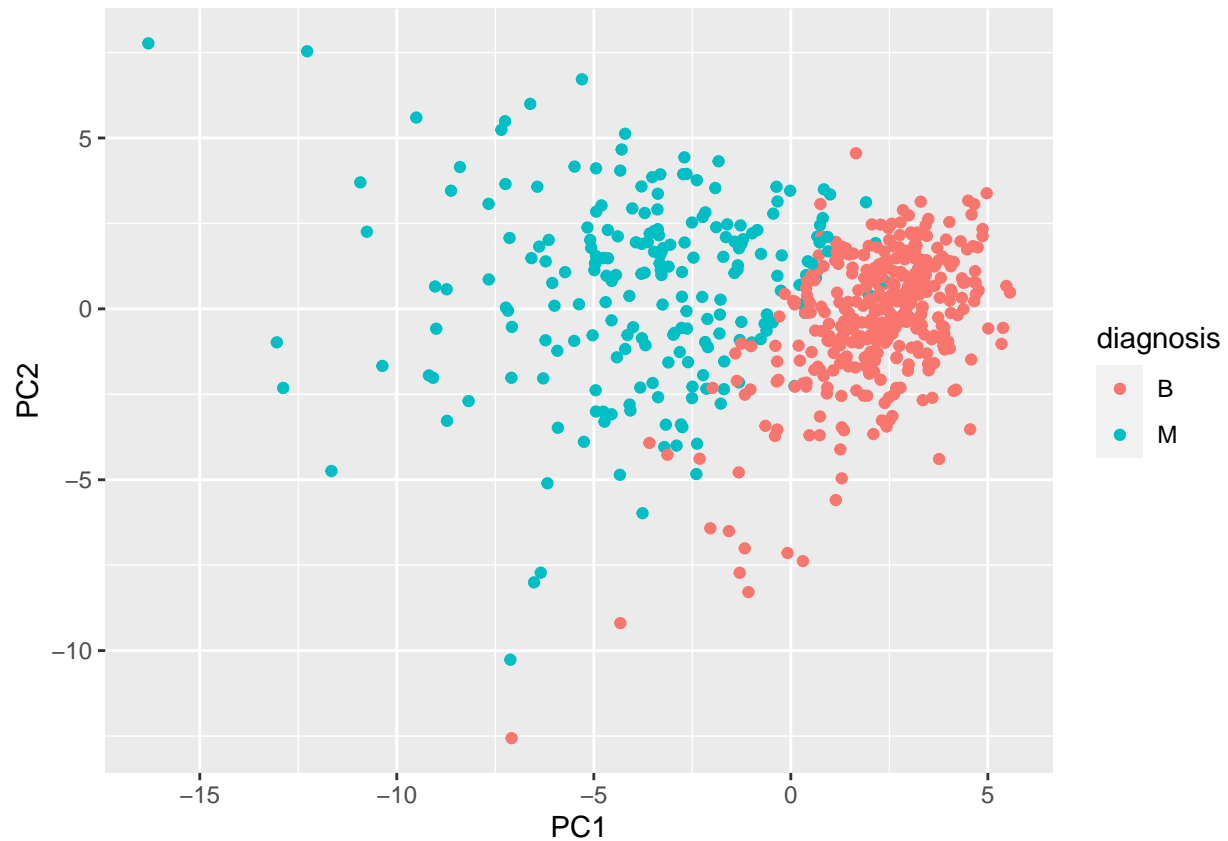
```
plot(wisc.pr$x[,c(1, 3)], col = diagnosis, xlab = "PC1", ylab = "PC3")
```



```
# Create a data.frame for ggplot
df <- as.data.frame(wisc.pr$x)
df$diagnosis <- diagnosis

# Load the ggplot2 package
library(ggplot2)

# Make a scatter plot colored by diagnosis
ggplot(df) +
  aes(PC1, PC2, col=diagnosis) +
  geom_point()
```



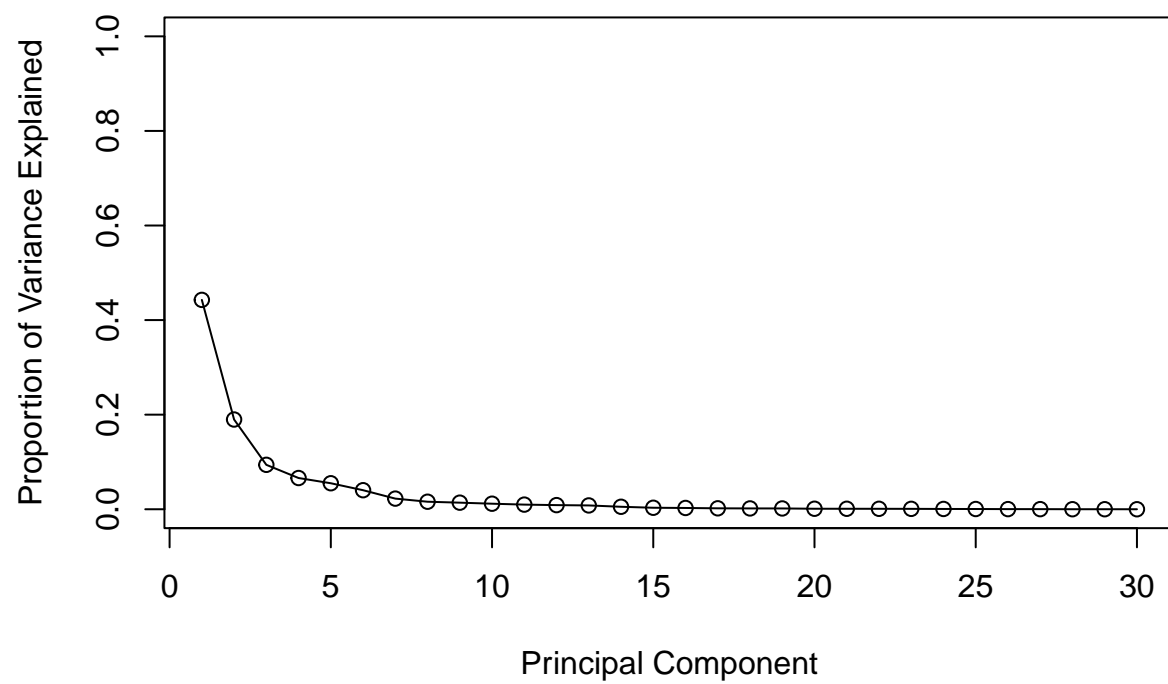
```
# Calculate variance of each component
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

```
## [1] 13.281608  5.691355  2.817949  1.980640  1.648731  1.207357
```


```
# Variance explained by each principal component: pve
pve <- pr.var/sum(pr.var)
```

```
# Plot variance explained for each principal component
```

```
plot(pve, xlab = "Principal Component", ylab = "Proportion of Variance Explained", ylim = c(0, 1), type = "p")
```



```
# Alternative scree plot of the same data, note data driven y-axis
barplot(pve, ylab = "Precent of Variance Explained", names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`?

```
wisc.pr$rotation[,1]["concave.points_mean"]
```

```
## concave.points_mean
## -0.2608538
```

Q10. What is the minimum number of principal components required to explain 80% of the variance of the data?

```
summary(wisc.pr)
```

```
## Importance of components:
##          PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  3.6444  2.3857  1.67867  1.40735  1.28403  1.09880  0.82172
## Proportion of Variance 0.4427  0.1897  0.09393  0.06602  0.05496  0.04025  0.02251
## Cumulative Proportion 0.4427  0.6324  0.72636  0.79239  0.84734  0.88759  0.91010
##          PC8      PC9      PC10     PC11     PC12     PC13     PC14
## Standard deviation  0.69037  0.6457  0.59219  0.5421  0.51104  0.49128  0.39624
## Proportion of Variance 0.01589  0.0139  0.01169  0.0098  0.00871  0.00805  0.00523
## Cumulative Proportion 0.92598  0.9399  0.95157  0.9614  0.97007  0.97812  0.98335
##          PC15     PC16     PC17     PC18     PC19     PC20     PC21
## Standard deviation  0.30681  0.28260  0.24372  0.22939  0.22244  0.17652  0.1731
```

```
## Proportion of Variance 0.00314 0.00266 0.00198 0.00175 0.00165 0.00104 0.0010
## Cumulative Proportion 0.98649 0.98915 0.99113 0.99288 0.99453 0.99557 0.9966
##          PC22      PC23      PC24      PC25      PC26      PC27      PC28
## Standard deviation    0.16565 0.15602 0.1344 0.12442 0.09043 0.08307 0.03987
## Proportion of Variance 0.00091 0.00081 0.0006 0.00052 0.00027 0.00023 0.00005
## Cumulative Proportion 0.99749 0.99830 0.9989 0.99942 0.99969 0.99992 0.99997
##          PC29      PC30
## Standard deviation    0.02736 0.01153
## Proportion of Variance 0.00002 0.00000
## Cumulative Proportion 1.00000 1.00000
```

We need at least 5 of principal components to explain 80% of the variance of the data.

3. Hierarchical clustering

```
# Scale the wisc.data data using the "scale()" function
data.scaled <- scale(wisc.data)
```


```
data.dist <- dist(data.scaled)
```

```
wisc.hclust <- hclust(data.dist, method="complete")
```

Q11. Using the `plot()` and `abline()` functions, what is the height at which the clustering model has 4 clusters?

```
plot(wisc.hclust)
abline(h=c(19, 20), col="red", lty=2)
```

Cluster Dendrogram



```
data.dist
hclust(*, "complete")
```

The clustering model has 4 clusters at the heights of 19 and 20.

```
wisc.hclust.clusters <- cutree(wisc.hclust, k=4)
```

```
table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters  B  M
##                   1 12 165
##                   2  2   5
##                   3 343  40
##                   4  0   2
```

Q12. Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10?

```
for (i in 2:10) {
  clusters <- cutree(wisc.hclust, k=i)
  print(table(clusters, diagnosis))
}
```

```
##              diagnosis
## clusters      B  M
##          1 357 210
##          2   0   2
```

```

##      diagnosis
## clusters  B  M
##      1 355 205
##      2   2   5
##      3   0   2
##      diagnosis
## clusters  B  M
##      1  12 165
##      2   2   5
##      3 343  40
##      4   0   2
##      diagnosis
## clusters  B  M
##      1  12 165
##      2   0   5
##      3 343  40
##      4   2   0
##      5   0   2
##      diagnosis
## clusters  B  M
##      1  12 165
##      2   0   5
##      3 331  39
##      4   2   0
##      5  12   1
##      6   0   2
##      diagnosis
## clusters  B  M
##      1  12 165
##      2   0   3
##      3 331  39
##      4   2   0
##      5  12   1
##      6   0   2
##      7   0   2
##      diagnosis
## clusters  B  M
##      1  12  86
##      2   0  79
##      3   0   3
##      4 331  39
##      5   2   0
##      6  12   1
##      7   0   2
##      8   0   2
##      diagnosis
## clusters  B  M
##      1  12  86
##      2   0  79
##      3   0   3
##      4 331  39
##      5   2   0
##      6  12   0
##      7   0   2


```

```
##      8  0  2
##      9  0  1
##      diagnosis
## clusters  B  M
##      1  12 86
##      2   0 59
##      3   0  3
##      4 331 39
##      5   0 20
##      6   2  0
##      7  12  0
##      8   0  2
##      9   0  2
##     10   0  1
```

I think that cutting into 4 clusters is the best in regard to diagnosis.


Q13. Which method gives your favorite results for the same data.dist dataset? Explain your reasoning.

```
wisc.hclust.complete <- hclust(data.dist, method="complete")
plot(wisc.hclust.complete)
```



```
wisc.hclust.single <- hclust(data.dist, method="single")  
plot(wisc.hclust.single)
```


Cluster Dendrogram



data.dist
hclust (*, "single")

```
wisc.hclust.average <- hclust(data.dist, method="average")  
plot(wisc.hclust.average)
```

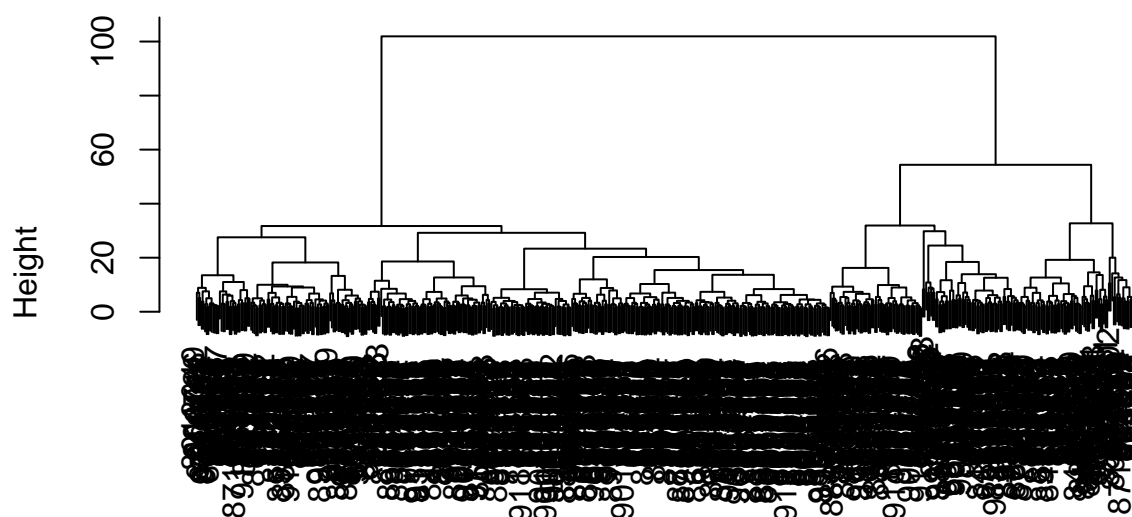
Cluster Dendrogram



data.dist
hclust (*, "average")

```
wisc.hclust.ward.D2 <- hclust(data.dist, method="ward.D2")  
plot(wisc.hclust.ward.D2)
```

Cluster Dendrogram



```
data.dist  
hclust (*, "ward.D2")
```

I think that “ward.D2” method gives me favorite results compared with others. I can get a good cluster vs diagnoses match only with the 2 clusters using this method.


4. OPTIONAL: K-means clustering

5. Combining methods

```
wisc.pr.hclust <- hclust(dist(wisc.pr$x[,1:4]), method="ward.D2")
```

```
plot(wisc.pr.hclust)  
abline(h=80, col="red")
```


Cluster Dendrogram



```
dist(wisc.pr$x[, 1:4])
hclust (*, "ward.D2")
```

Let's find out cluster membership vector by cutting this tree into k=2 group.

```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
## grps
## 1 2
## 171 398
```

Now let's compare to the expert M and B vector.

```
table(diagnosis)
```


```
## diagnosis
## B M
## 357 212
```

We can do a cross-table by giving the 'table()' function two inputs.

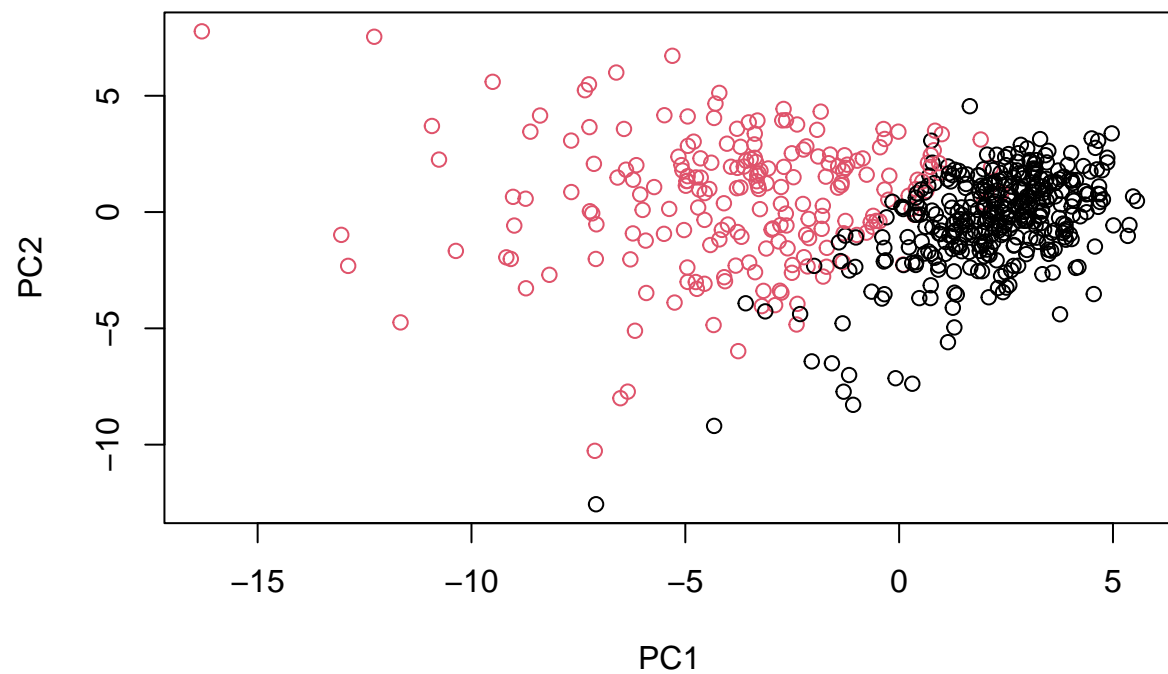
```
table(grps, diagnosis)
```

```
## diagnosis
## grps B M
## 1 6 165
## 2 351 47
```

```
plot(wisc.pr$x[,1:2], col=grps)
```



```
plot(wisc.pr$x[,1:2], col=diagnosis)
```




```
g <- as.factor(grps)
levels(g)
```

```
## [1] "1" "2"
```

```
g <- relevel(g,2)
levels(g)
```

```
## [1] "2" "1"
```

```
plot(wisc.pr$x[,1:2], col=g)
```



Q15. How well does the newly created model with four clusters separate out the two diagnoses?

```
table(grps, diagnosis)
```

```
##      diagnosis
## grps   B    M
##    1    6 165
##    2   351  47
```

Q16. How well do the k-means and hierarchical clustering models you created in previous sections (i.e. before PCA) do in terms of separating the diagnoses? Again, use the table() function to compare the output of each model (wisc.km\$cluster and wisc.hclust.clusters) with the vector containing the actual diagnoses.

```
table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters   B    M
##              1   12 165
##              2    2   5
##              3  343  40
##              4    0   2
```

(As section 4 is optional, I do not have the output of the k-means model)

6. Sensitivity/Specificity

Accuracy, essentially how many did we get correct?

```
(165+351)/nrow(wisc.data)
```

```
## [1] 0.9068541
```

Q17. Which of your analysis procedures resulted in a clustering model with the best specificity?
How about sensitivity?

Sensitivity

```
(165)/(165+47)
```

```
## [1] 0.7783019
```

Specificity

```
(351)/(351+47)
```

```
## [1] 0.8819095
```


I just calculated the specificity and the sensitivity of the model with the “ward.D2” method.

7. Prediction

```
#url <- "new_samples.csv"
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```

```
##          PC1          PC2          PC3          PC4          PC5          PC6          PC7
## [1,]  2.576616 -3.135913  1.3990492 -0.7631950  2.781648 -0.8150185 -0.3959098
## [2,] -4.754928 -3.009033 -0.1660946 -0.6052952 -1.140698 -1.2189945  0.8193031
##          PC8          PC9          PC10          PC11          PC12          PC13          PC14
## [1,] -0.2307350 0.1029569 -0.9272861 0.3411457  0.375921 0.1610764 1.187882
## [2,] -0.3307423 0.5281896 -0.4855301 0.7173233 -1.185917 0.5893856 0.303029
##          PC15          PC16          PC17          PC18          PC19          PC20
## [1,] 0.3216974 -0.1743616 -0.07875393 -0.11207028 -0.08802955 -0.2495216
## [2,] 0.1299153 0.1448061 -0.40509706 0.06565549 0.25591230 -0.4289500
##          PC21          PC22          PC23          PC24          PC25          PC26
## [1,] 0.1228233 0.09358453 0.08347651 0.1223396 0.02124121 0.078884581
## [2,] -0.1224776 0.01732146 0.06316631 -0.2338618 -0.20755948 -0.009833238
##          PC27          PC28          PC29          PC30
## [1,] 0.220199544 -0.02946023 -0.015620933 0.005269029
## [2,] -0.001134152 0.09638361 0.002795349 -0.019015820
```

```
plot(wisc.pr$x[,1:2], col=g)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



Q18. Which of these new patients should we prioritize for follow up based on your results?

Based on the results, the patients in the group 2 should be prioritized for follow up.